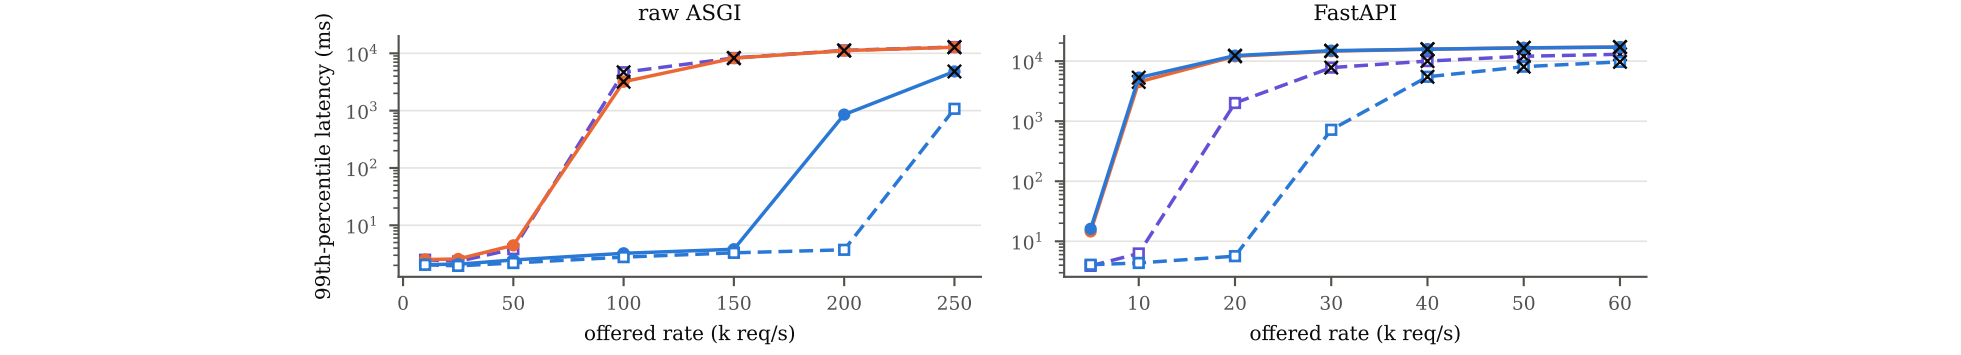
raw ASGI
FastAPI
99th-percentile latency (ms)
offered rate (k req/s)
offered rate (k req/s)
Uvicorn (httptools+uvloop), processes
uringpy-http, GIL per request, threads
uringpy-http, GIL per batch, threads
uringpy-http, GIL per batch, processes
delivered < 95% of offered